

Performance Analysis and Optimization

Giannis Koutsou
LAP2015, Oct. 14-16, Jülich



THE CYPRUS
INSTITUTE

Resources

- **Reference card for OpenMP**

- <http://openmp.org/mp-documents/OpenMP3.1-CCard.pdf>

- **Reference for Vector-Intrinsics**

- <https://software.intel.com/sites/landingpage/IntrinsicsGuide/>

- **Exercises**

- github.com/g-koutsou/LAP2015/

```
git clone https://github.com/g-koutsou/LAP2015.git
```

Exercises

■ Generic instructions:

- A Makefile for each exercise subdirectory
 - Type `make` to compile
- A job script, `run.sh`, requesting 1 node for 5-15 mins.
 - Type `qsub run.sh` to submit job
 - Type `qstat -u $USER` to see the status (queued, running, completing)
 - Upon completion a file with a `.log` extension contains the output of the run
- `__TODO__` tags in code indicate where you need to add code for completing the exercises

Exercise 1

■ Exercise 1

- Subdirs: Ex1a, Ex1b, Ex1c

■ Ex1a

- Complex, single precision, **axpy** kernel:

```
for(i=0; i<L; i++)  
    y[i] = a*x[i] + y[i]
```

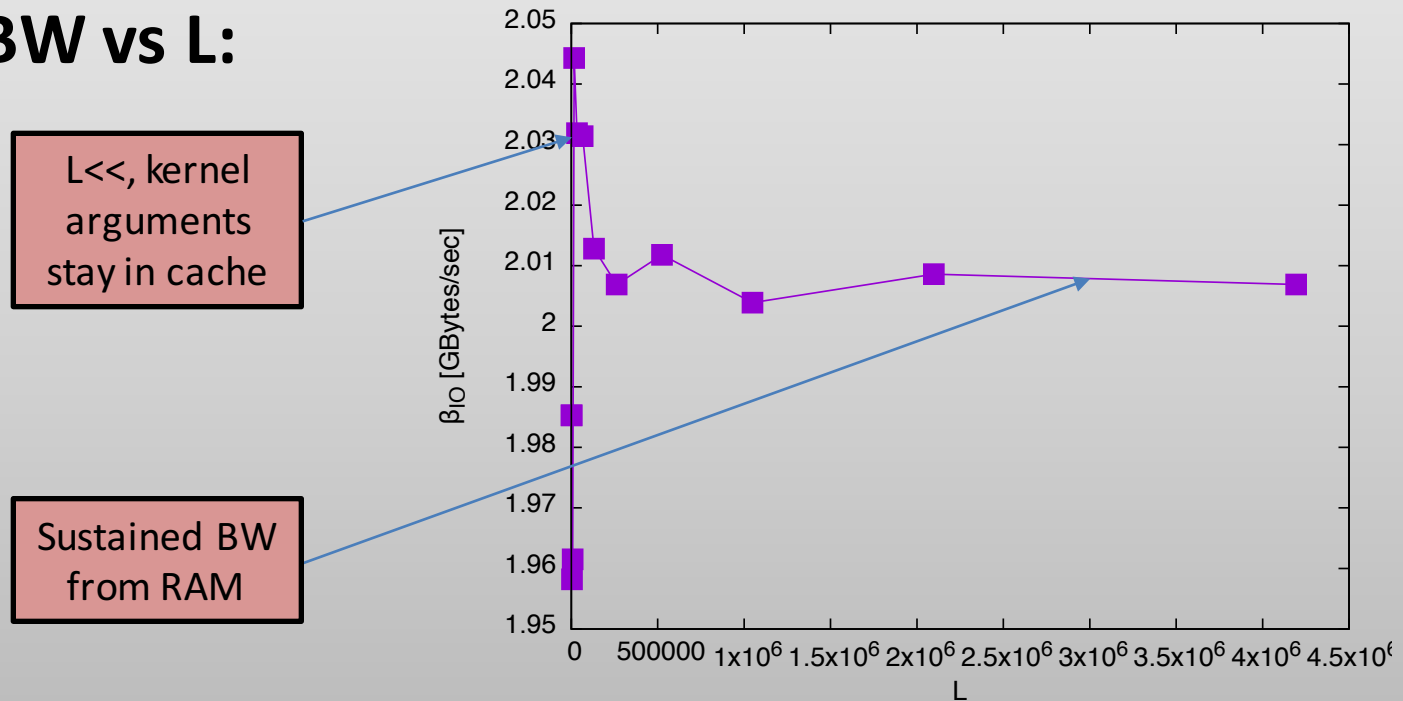
- Modify program to report
 - Time per kernel call (μsec)
 - Floating point rate (Gflop/sec)
 - I/O rate (Gbytes/sec)

Exercise 1a

■ For example:

$L = 1048576$, $1.26e+04 \pm 3.53e+00$ usec/call, perf. = 0.668 GFlop/sec, bw = 2.004 GBytes/sec

■ Plot BW vs L:



Exercise 1b

■ Implement Kernel using vector-intrinsics

```
for(i=0; i<L; i+=2) {  
    y[i].re = a.re*x[i].re - a.im*x[i].im + y[i].re  
    y[i].im = a.re*x[i].im + a.im*x[i].re + y[i].im  
    y[i+1].re = a.re*x[i+1].re - a.im*x[i+1].im + y[i+1].re  
    y[i+1].im = a.re*x[i+1].im + a.im*x[i+1].re + y[i+1].im  
}
```

Exercise 1b

■ Implement Kernel using vector-intrinsics

```
for(i=0; i<l; i+=2) {  
    y[i].re = a.re*x[i].re - a.im*x[i].im + y[i].re  
    y[i].im = a.re*x[i].im + a.im*x[i].re + y[i].im  
    y[i+1].re = a.re*x[i+1].re - a.im*x[i+1].im + y[i+1].re  
    y[i+1].im = a.re*x[i+1].im + a.im*x[i+1].re + y[i+1].im  
}
```

Exercise 1b

■ Implement Kernel using vector-intrinsics

```
for(i=0; i<l; i+=2) {  
    y[i].re = a.re*x[i].re - a.im*x[i].im + y[i].re  
    y[i].im = a.re*x[i].im + a.im*x[i].re + y[i].im  
    y[i+1].re = a.re*x[i+1].re - a.im*x[i+1].im + y[i+1].re  
    y[i+1].im = a.re*x[i+1].im + a.im*x[i+1].re + y[i+1].im  
}
```


Exercise 1b

■ Implement Kernel using vector-intrinsics

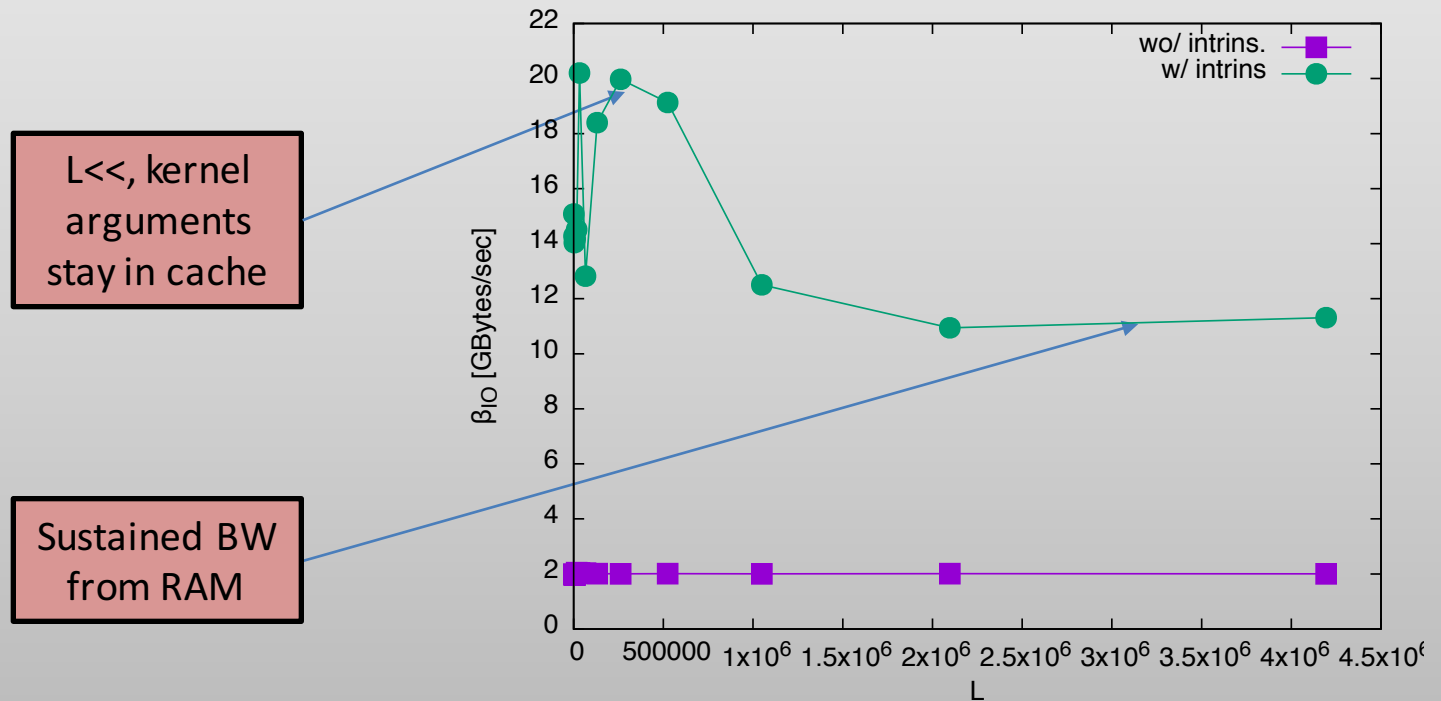
```
for(i=0; i<L; i+=2) {  
    y[i].re = a.re*x[i].re - a.im*x[i].im + y[i].re  
    y[i].im = a.re*x[i].im + a.im*x[i].re + y[i].im  
    y[i+1].re = a.re*x[i+1].re - a.im*x[i+1].im + y[i+1].re  
    y[i+1].im = a.re*x[i+1].im + a.im*x[i+1].re + y[i+1].im  
}
```

■ For every (double-) iteration

- One shuffle
- Two add
- Two mul
- Two loads, one store

Exercise 1b

■ Plot BW vs L:

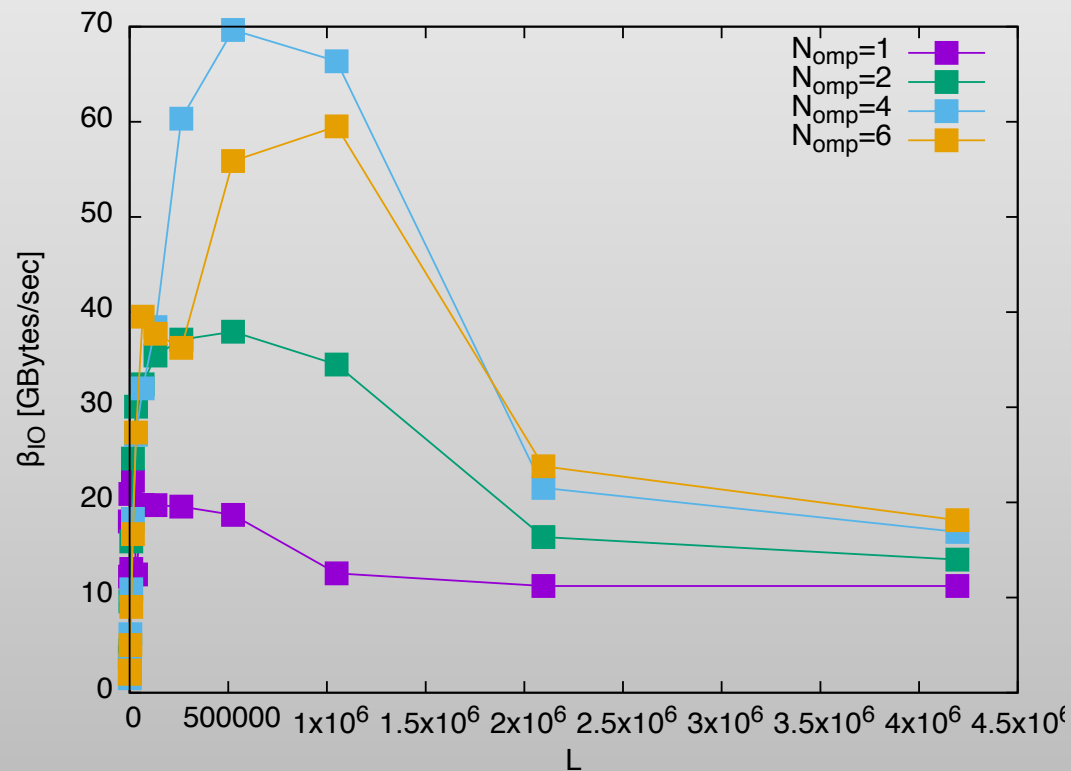


Exercise 1c

- **Add OMP pragmas to parallelize the i-loop**
- **Use:**
 - `#pragma omp parallel`
 {
 ...
 }
 - `#pragma omp for`
- **Which variables need to be made local (private)?**

Exercise 1c

- Plot BW vs L, for 1, 2, 4, and 6 OMP threads:



Exercise 2

■ Ex2

- Array of double NxN matrices times constant NxN matrix,

mxam:

```
for(i=0; i<L; i++)  
  for(a=0; a<N; a++)  
    for(b=0; b<N; b++)  
      for(c=0; c<N; c++)  
        y[i][a][b] += A[a][c]*x[i][c][b]
```

- Purpose is **not** to optimise this kernel
- This kernel will allow us to explore the transition between computational-bound and bandwidth-bound kernels (how?)

Exercise 2

■ Ex2

- Array of double NxN matrices times constant NxN matrix,

mxam:


```
for(i=0; i<L; i++)  
  for(a=0; a<N; a++)  
    for(b=0; b<N; b++)  
      for(c=0; c<N; c++)  
        y[i][a][b] += A[a][c]*x[i][c][b]
```

- As in Ex1, modify program to report
 - Time per kernel call (μsec)
 - Floating point rate (Gflop/sec)
 - I/O rate (Gbytes/sec)
- Run and plot: $T/(N**2*L)$ vs N

Exercise 2

■ Ex2

— Run and plot $T/(N^2L)$ vs N



$$T_{FP} = N^2 L (2N - 1) \frac{1}{P \beta_{FP}}$$

$$T_{IO} = N^2 L \frac{1}{\beta_{IO}(P)}$$

$$\frac{T}{N^2 L} = \max\left[\frac{2N - 1}{P \beta_{FP}}, \frac{1}{\beta_{IO}(P)}\right]$$